

Yash Arvind

Madison, WI · yasharvind12345@gmail.com · 608-515-1932 · linkedin.com/in/yash-arvind-294516218

Open to relocation — Mountain View, CA · New York, NY · San Bruno, CA · San Jose, CA

SUMMARY

Technical product builder who turns ambiguous financial, billing, and AI workflows into shipped products with measurable outcomes. Owned vision, discovery, requirements, prioritization, and rollout at Zendesk, closing a five-year data-integrity gap and saving **9,000+ analyst hours** and **\$2-3M** annually. Applies AI/ML end to end — agent orchestration, RAG, NLP extraction, anomaly detection — with the fluency to judge where models help and where deterministic systems win. Founder of three ventures taken from 0 to 1.

EDUCATION

University of Wisconsin-Madison

Expected May 2027 · Madison, WI

B.S. Data Science & Economics · GPA: 3.4

Coursework: Deep Learning & Generative Models, AI/ML, Database Systems (SQL), Algorithms, Linear Algebra, Statistical Computing (R)

EXPERIENCE

Zendesk — AI Product Intern, Business Operations

Jun 2026 – Aug 2026

- Drove an AI-assisted Salesforce CPQ and Zuora reconciliation product from problem framing to production, converting a **\$20M** cross-system discrepancy into auditable matching logic — closing a five-year data-integrity gap and saving **9,000+ analyst hours** and **\$2-3M** annually.
- Architected the core matching layer in deterministic Snowflake SQL and reserved AI for the surrounding classification and simulation layers, keeping output auditable while validating on 100K+ records at **93%+ accuracy**.
- Worked across Finance, RevOps, and Engineering to translate contract structures — proration, tiered pricing, discounting, ramps, and ELAs — into calculation logic, system requirements, and API specifications.
- Partnered with VP and SVP stakeholders to define product scope, phasing ramp/ELA contracts into a follow-on release; tracked requirements and releases in Jira and Confluence.
- Built a MiroFish agent-simulation environment for extensive UAT, spawning **millions of AI agents** with varied characteristics to model real-world conditions and human behavior.
- Shipped an LLM token-optimization tool that cut internal AI spend **~20%**.
- Built a retrieval-based documentation-freshness system flagging conflicting content across Help Center, Google Drive, and Unleash, with Claude drafting corrections.

Ernst & Young (EY) — M&A Consulting Intern

Jun 2025 – Aug 2025

- Performed financial due diligence on a live M&A deal, building DCF valuation models in Excel (Pivot Tables, XLOOKUP, SUMIFS, INDEX-MATCH, Power Query) to support buy-side decisions.
- Reviewed **200+ pages** of target financials to produce an Independent Accountants' Report, synthesizing findings into executive-ready briefs for senior stakeholders and clients.
- Produced red-flag reports surfacing financial risks across the target's statements and built Power BI dashboards to present diligence findings to the engagement team.

Synechron Technologies — AI Engineering Intern

Jul 2023 – Aug 2023

- Partnered with claims subject-matter experts to translate valuation logic into requirements, trained an ML cost estimator (scikit-learn, GridSearchCV) on 10,000+ claims records, and served it behind a REST API — lifting accuracy **34%**, cutting processing time **70%**.

PROJECTS

TAM (Transaction Analysis Machine) — Co-Founder & CEO

2025 – Present

- Set product vision and roadmap for an agentic due-diligence platform: LangGraph orchestration over a ChromaDB RAG layer and the Claude API, with a LangChain + FinBERT pipeline parsing SEC filings and contracts for covenants, risk clauses, and filing sentiment.
- Layered Isolation Forest anomaly detection over financial statements to flag irregularities, auto-generating PDF/Excel reports in **under 8s** per company on a FastAPI/Celery, PostgreSQL, and React stack.
- Run discovery with due-diligence practitioners and Managing Directors at EY and PwC, converting customer feedback into prioritized requirements and a go-to-market plan.

Cortexa — Co-Founder

2025 – Present

- Own product strategy for an enterprise LLM hallucination-detection layer using token-level entropy and retrieval-consistency checks — targeting **60%+ error reduction** in RAG pipelines and ~30% token cost reduction, validated with Nvidia and AmFam.
- Designed it to drop into existing LangChain and LlamaIndex retrieval stacks at under 12ms overhead.

Flux — Co-Founder

2025 – Present

- Shipped a live iOS co-founder matching app to the App Store, scaling to **350 users**.

EarningLens — 1st Place, Cursor Hackathon

2025

- Real-time earnings-call analyzer ingesting live transcripts via the SEC EDGAR API, running sentiment and NER extraction so users can fact-check the guidance and claims executives make on the call within **3s**.

SKILLS

Product Management: Product discovery · Requirements definition · Prioritization & scope tradeoffs · Roadmapping · Cross-functional collaboration · Stakeholder management · GTM strategy

AI & ML: Claude API · LangChain/LangGraph · RAG · agent orchestration · tool use · prompt & eval design · NLP/NER · FinBERT · spaCy · PyTorch · scikit-learn

Engineering & Data: Python · SQL · R · FastAPI · REST APIs · PostgreSQL · ChromaDB · Snowflake · Redis/Celery · Docker · Git

Systems & Analytics: Jira (Atlassian) · Confluence · Salesforce CRM/CPQ · Zuora · Tableau · Power BI · Advanced Excel · Financial modeling · DCF valuation